

HW5 Report

Part 1: Site Design and README

Project Structure

```
CS6675_HW5/
├─ server.js      # express server with routing
├─ package.json   # node dependencies
└─ public/        # site pages
    ├─ index.html
    ├─ page1.html
    ├─ page2.html
    ├─ page3.html
    ├─ page4.html
    ├─ page5.html
    └─ images/
```

Features

- Hosts a very simple web server with some additional QoL packages
 - **Compression:** Uses gzip compression for better performance
 - **Logging:** Morgan middleware for HTTP request logging

Dependencies

- **express:** Web framework for Node.js
- **compression:** Gzip compression middleware
- **morgan:** HTTP request logger
- **apicache:** Caching middleware (commented out)

Installation

* This is only necessary if cloning from git, if using zip skip to **usage**

1. Clone or download the project
2. Install dependencies:

```
npm install
```

Usage

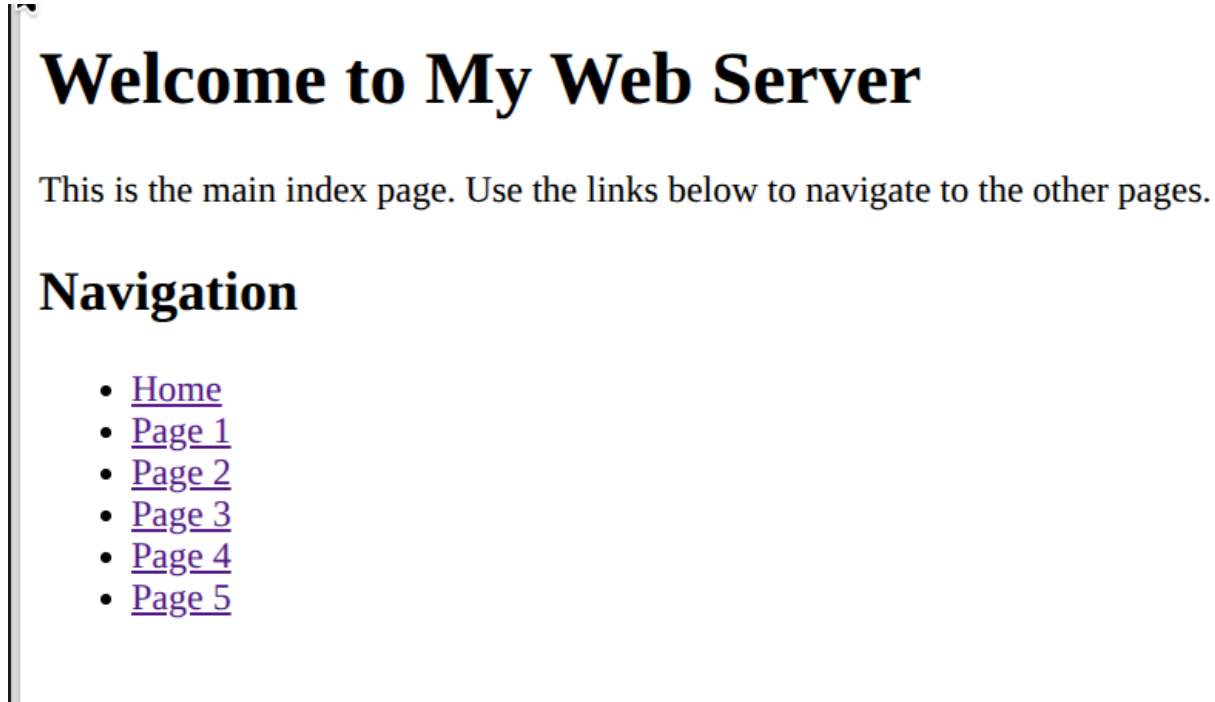
1. Start the server:

```
node server.js
```

2. Open your browser and navigate to:

```
http://localhost:3000
```

3. The server will be accessible at port 3000 and will display:



Click through the pages as needed

Server Configuration

This is easily changed in `server.js`

- **Port:** 3000
- **Host:** 0.0.0.0 (accessible from all network interfaces)

Optional Features

The server includes code for **caching** using apicache middleware. This **significantly improves performance**. To enable these features, uncomment the lines in `server.js`

For performance testing instructions see section 3

Part 2: Remote Client Access

I set up remote client access by utilizing tailscale. Tailscale is a network overlay that sits on top of any existing network infrastructure creating a private network allowing devices on the tailscale network to freely discover each other as if they were on the same local network (essentially acting like a multi-device VPN). This solution was completely new to me to set up and get running but it worked quite well allowing me to access my localhost website on any network (including mobile data as seen in the mobile screenshot) **without** using a hotspot or a

shared physical network. One positive consequence of setting up a tailscale network was that it let me freely share files across the "local" network allowing me to get these screenshots directly to my computer without any complicated sharing methods (AirDrop, cloud sync, etc.).

```
Unknown adapter Tailscale:

Connection-specific DNS Suffix . : tail29df4c.ts.net
IPv6 Address. . . . . : fd7a:115c:a1e0::101:1134
Link-local IPv6 Address . . . . . : fe80::51db:8521:5771:6e3b%26
IPv4 Address. . . . . : 100.81.17.52
Subnet Mask . . . . . : 255.255.255.255
Default Gateway . . . . . :

Wireless LAN adapter Local Area Connection* 1:

Media State . . . . . : Media disconnected
Connection-specific DNS Suffix . :

Wireless LAN adapter Local Area Connection* 2:

Media State . . . . . : Media disconnected
Connection-specific DNS Suffix . :

Wireless LAN adapter Wi-Fi:

Connection-specific DNS Suffix . :
Link-local IPv6 Address . . . . . : fe80::a6d7:abcc:3259:da71%11
IPv4 Address. . . . . : 10.146.35.107
Subnet Mask . . . . . : 255.0.0.0
Default Gateway . . . . . : 10.128.128.128

Ethernet adapter Bluetooth Network Connection:

Media State . . . . . : Media disconnected
Connection-specific DNS Suffix . :
```

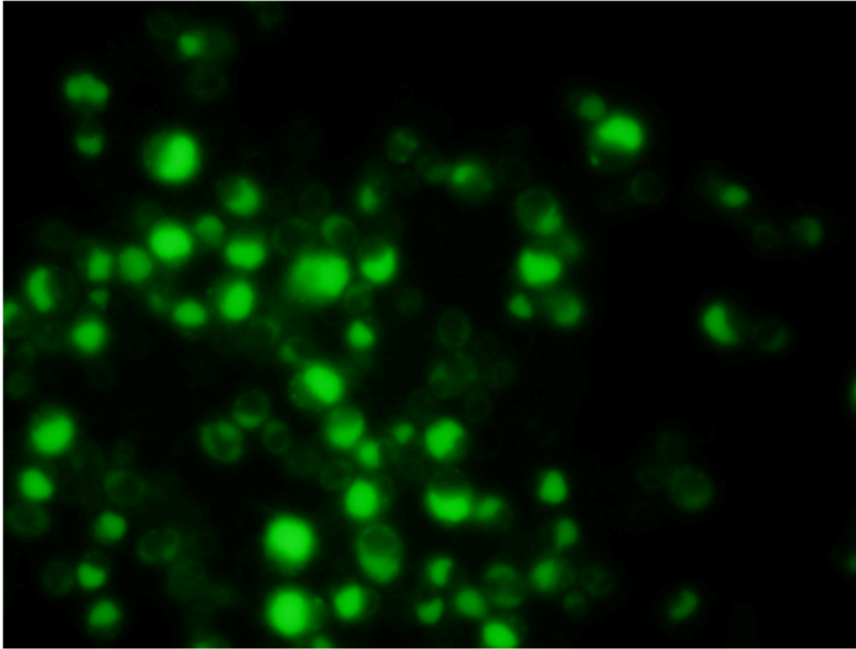
IP-Config on Windows - See Tailscale at the top



[Home](#) | [Page 1](#) | [Page 2](#) | [Page 3](#) | [Page 4](#) | [Page 5](#)

Three Image Layout

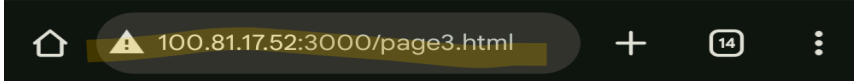
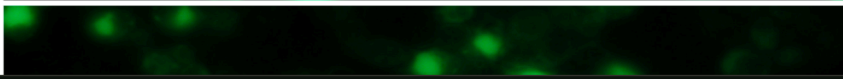
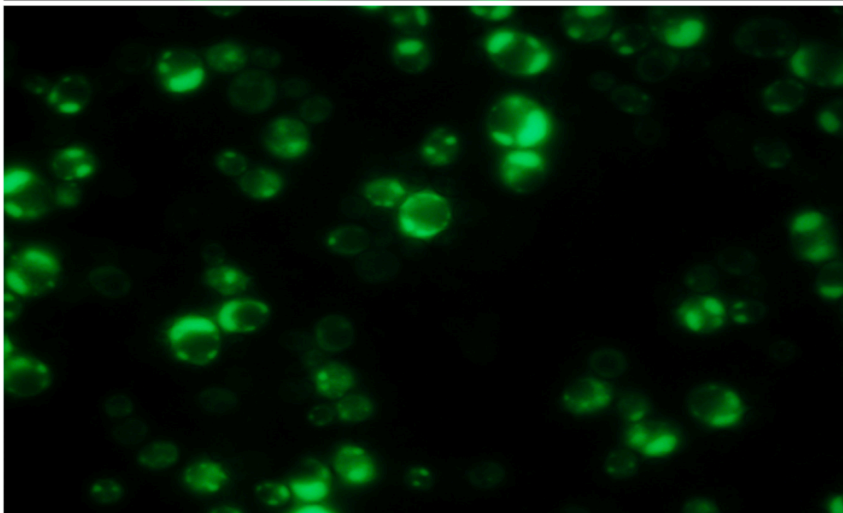
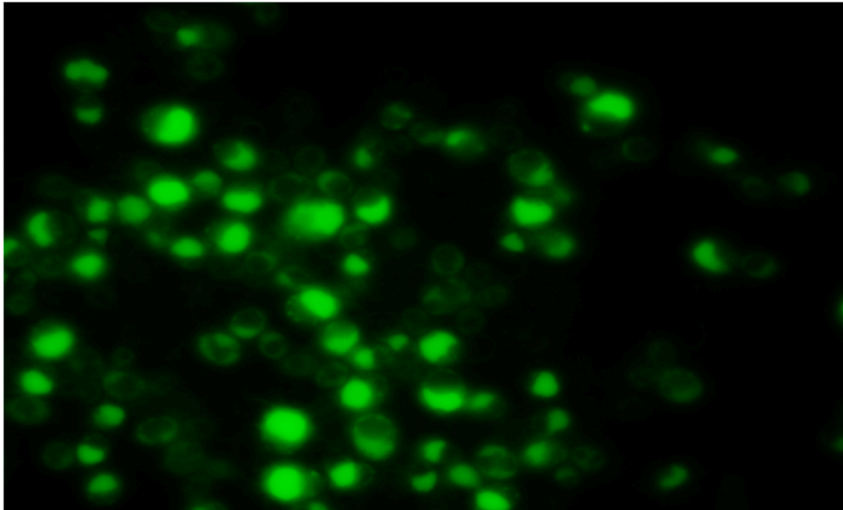
This page displays three different images.



Localhost site being hosted via Express and Node

Three Image Layout

This page displays three different images.



Website being accessible on my android phone

As you might notice the IP address here is different from the one in my `ipconfig` screenshot, this is because I had to switch to my Linux machine as I was running into some stability issues with windows and it had better compatibility with the benchmarking software I was using. Tailscale allowed it so that I could still link all of my devices together and get it up and running quite quickly.

Part 3 and 4: Stress Test and Latency/Throughput Analysis

The tool I used to test the performance of my web-server was [called oha](#), a rust-based TUI that was modern, well-threaded, and gave great data-visibility compared to the gold-standard **apachebench**.

```

Summary:
Success rate: 100.00%
Total: 30001.6189 ms
Slowest: 80.3855 ms
Fastest: 2.3995 ms
Average: 8.1403 ms
Requests/sec: 6142.0019

Total data: 153.37 MiB
Size/request: 873 B
Size/sec: 5.11 MiB

Response time histogram:
2.399 ms [1] |
10.198 ms [170760] |
17.997 ms [12709] |
25.795 ms [658] |
33.594 ms [71] |
41.392 ms [5] |
49.191 ms [3] |
56.990 ms [4] |
64.788 ms [3] |
72.587 ms [3] |
80.385 ms [3] |

Response time distribution:
10.00% in 6.8976 ms
25.00% in 7.1201 ms
50.00% in 7.5601 ms
75.00% in 8.7873 ms
90.00% in 9.7304 ms
95.00% in 10.7960 ms
99.00% in 14.8846 ms
99.90% in 24.3072 ms
99.99% in 37.2399 ms

```



Some examples of the outputs of `oha`, whereas `apachebench` simply outputs a couple of text lines of data.

The command used to run the following tests was as follows:

```
oha -z s -c r http://localhost:3000/page1.html
```

where

- -z s is how many seconds to run the test (I used 30 seconds)
 - This can be replaced by -n for number of requests
- -c r is how many concurrent requests per second (this was modified)

Performance Results (No Caching)

Test Run	Concurrent Users	Requests/sec (Throughput)	Avg Latency (ms)	Fastest (ms)	Slowest (ms)
Latency Test	10	5,932	1.69	0.85	16.20
Latency Test	50	6,142	8.14	2.40	80.39
Stress Test	200	5,935	33.71	20.02	721.40
Stress Test (Re-run)	400	5,906	67.70	45.00	2,588.70
Stress Test	1000	5,527	181.00	47.00	1,703.40

Without caching, the server handled a max of 6,142 req/sec, with it slowly leveling off as concurrent users grew, showing that the server slowly became overloaded. We also see this with the dramatic jump in latency as we 100x users our latency grows accordingly. Looking at the Response Time Percentiles chart as well we see that the range in response times grows quite significantly.

Performance Results (Caching)

Test Run	Concurrent Users	Requests/sec (Throughput)	Avg Latency (ms)	Fastest (ms)	Slowest (ms)
Latency Test	10	8,401	1.20	0.10	24.73
Latency Test	50	8,250	6.06	0.15	246.89
Stress Test	200	7,868	25.40	4.60	2,044.20
Stress Test	400	7,850	50.90	1.50	5,787.10
Stress Test	1000	7,955	106.00	26.00	2,984.00

With caching, the server stores copies of pages and is able to serve it instantly instead of rebuilding on every request. With caching we see a significant jump in throughput (an almost ~37% increase when comparing peaks). While the pattern for throughput and latency is similar to our data without caching (suggesting that caching gets to a point of diminishing returns), we see that the average latency across the board is lower (~41% decrease when comparing worst latency). This shows that caching is a resounding success especially when we look at the median response times which show an even greater improvement than when looking at averages that take into account the large outliers.

Response Time Percentiles (ms)

Test Run	Concurrent Users	P50 (Median)	P75	P90	P95	P99
Latency Test	10	1.51	1.61	2.01	2.99	3.92
Latency Test	50	7.56	8.79	9.73	10.80	14.88
Stress Test	200	32.67	34.81	37.52	40.60	50.78
Stress Test (Re-run)	400	65.50	68.50	74.50	79.30	90.20
Stress Test	1000	170.00	180.00	191.00	199.00	206.00

Response Time Percentiles (ms) - Caching

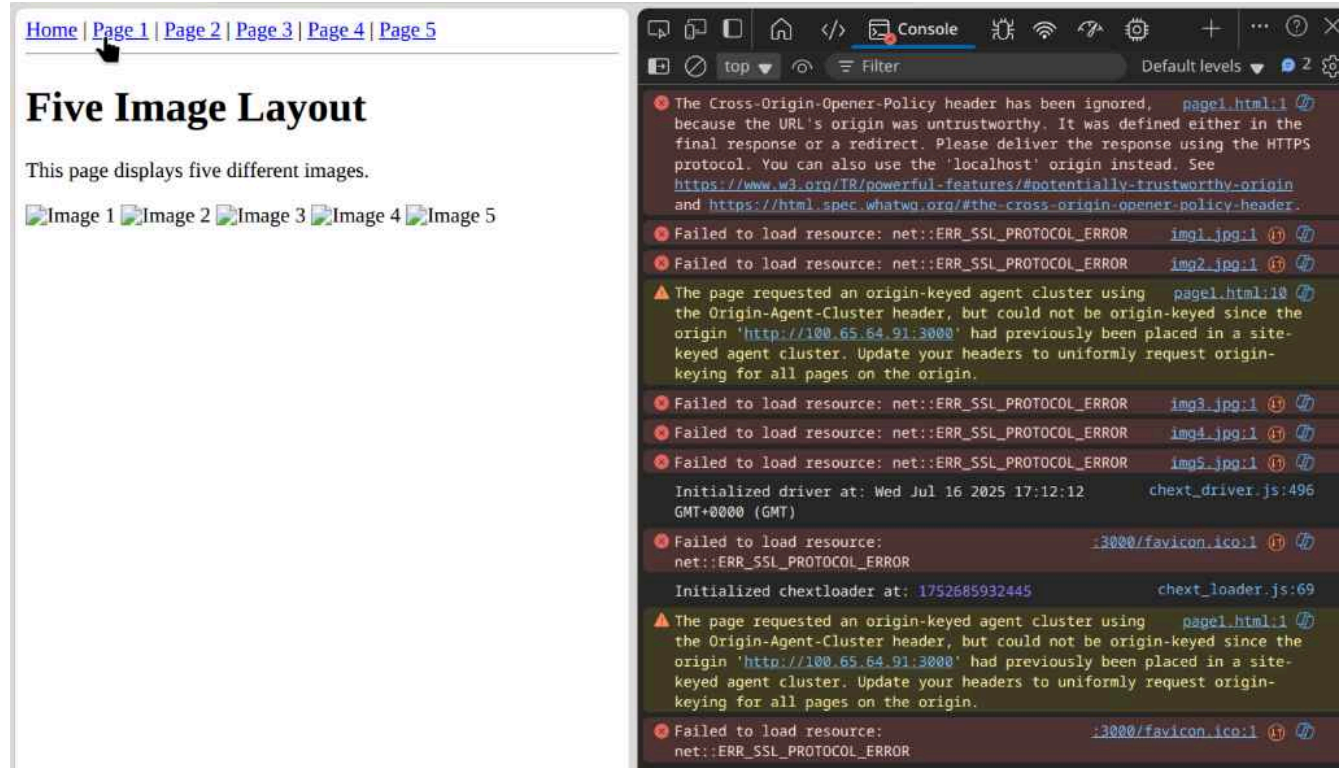
Test Run	Concurrent Users	P50 (Median)	P75	P90	P95	P99
Latency Test	10	1.03	1.09	1.40	2.40	3.35
Latency Test	50	5.40	6.76	7.71	8.65	11.53
Stress Test	200	24.00	25.80	30.10	32.70	38.40
Stress Test	400	48.20	50.90	56.20	61.60	78.90
Stress Test	1000	85.00	96.00	103.00	109.00	123.00

Experience

Overall, I thought this was a really cool project, as I pushed myself to learn a variety of new toolkits and software. Coming into this I had some familiarity with Apache TomCat so I wanted to explore newer more modern frameworks. I decided on using `Express.js` and `Node.js` because both are relatively modern and they are becoming increasingly popular amongst web developers today, it is well documented and very easy and quick to get up and running. In addition, it is very flexible allowing me to work with and explore different packages (like `morgan` and `compression`), which provided me with greater logging visibility and improved web-server performance right off the bat. Building the web server itself is where I ran into a couple of learnings:

1. Browser security policies have evolved a lot. Setting up the server I ran into numerous issues with multiple different browsers forcing my site to run with HTTPS even if the main site loaded in via HTTP. This is where I learned about HSTS or HTTP Strict Transport Security which tells browsers to always use HTTPS which while **not** configured for localhost was set to automatically work on my tailscale IP addresses which lead to quite a headache trying to figure out why my http requests were being sent as HTTPS and promptly

rejected.



2. The second thing I had to learn was how to actually access my webserver. Normally I would just find the IP of the device hosting the site and connect to it (say at my home router), but this failed to work in my apartment and on campus leading to a lot of confusion as to why. After doing some digging I found that this was a security feature that leads to device discover-ability being disabled. I searched for solutions like setting up a hotspot, which worked but ate up a lot of my bandwidth as I made the images quite large on my website. I also tried ZeroTier which allowed for the creation of a private network but was fairly complicated so eventually I settled on tailscale which had a very intuitive quick start guide and allowed me to get a private network up and running quite quickly.

In the theme of using modern tools, I wanted to use a different tool than apachebench which came preinstalled with my fedora install and is well-known for being quick and easy to use. The issue was that the application is famously singlethreaded which could lead to performance issues during benchmarking, so I looked for packages built on a more modern framework and came across `oha` which is a rust based benchmarking tool that had an excellent TUI that gave quite a bit of insight on what was happening at a network level and how the requests were being handled on my webserver. This combined with the histograms for request data made analysis of the data easier than if I was to use apachebench. Which leads me to learning 3:

3. Logging is incredibly helpful, for example in the first learning about security. Web-server logging is quite complete and thorough from the server side of things to the browser side of things with developer tools and being able to inspect network queries. Developer tools combined with the logging package I was using (Morgan), allowed me to identify that a security package I was using called helmet was erroneously denying mass requests leading to performance numbers that didn't make sense at all. Eventually I realized that the requests were being denied and I have since removed the package. Logging also helped with `oha` as it allowed me to see exactly what was happening when I ran the stress test to make sure that requests were being properly handled, I ran into an issue where the cached requests were seemingly performing worse than the uncached requests and I narrowed it down to how long it would take for my cache to dump stale entries (it was far too short).

Appendix (Raw Results)

Part 3 without Caching

```
oha -z 30s -c 50 http://localhost:3000/page1.html (Without Caching)
```

Summary :

Success rate: 100.00%

Total: 30001.6189 ms

Slowest: 80.3855 ms

Fastest: 2.3995 ms

Average: 8.1403 ms

Requests/sec: 6142.0019

Total data: 153.37 MiB

Size/request: 873 B

Size/sec: 5.11 MiB

Response time histogram:

2.399 ms [1]

10.198 ms [170760]

17.997 ms [12709]

25.795 ms [658]

33.594 ms [71]

41.392 ms [5]

49.191 ms [3]

56.990 ms [4]

64.788 ms [3]

72.587 ms [3]

80.385 ms [3]

Response time distribution:

10.00% in 6.8976 ms

25.00% in 7.1201 ms

50.00% in 7.5601 ms

75.00% in 8.7873 ms

90.00% in 9.7304 ms

95.00% in 10.7960 ms

99.00% in 14.8846 ms

99.90% in 24.3072 ms

99.99% in 37.2399 ms



Visual Representation of the TUI

Summary:

Success rate: 100.00%
Total: 30001.6189 ms
Slowest: 80.3855 ms
Fastest: 2.3995 ms
Average: 8.1403 ms
Requests/sec: 6142.0019

Total data: 153.37 MiB
Size/request: 873 B
Size/sec: 5.11 MiB

Response time histogram:

2.399 ms	[1]	
10.198 ms	[170760]	
17.997 ms	[12709]	
25.795 ms	[658]	
33.594 ms	[71]	
41.392 ms	[5]	
49.191 ms	[3]	
56.990 ms	[4]	
64.788 ms	[3]	
72.587 ms	[3]	
80.385 ms	[3]	

Response time distribution:

10.00% in 6.8976 ms
25.00% in 7.1201 ms
50.00% in 7.5601 ms
75.00% in 8.7873 ms
90.00% in 9.7304 ms
95.00% in 10.7960 ms
99.00% in 14.8846 ms
99.90% in 24.3072 ms
99.99% in 37.2399 ms

```
Success rate: 100.00%
Total:      30001.1710 ms
Slowest:    16.1964 ms
Fastest:    0.8513 ms
Average:    1.6850 ms
Requests/sec: 5932.1018
```

```
Total data: 148.16 MiB
Size/request: 873 B
Size/sec: 4.94 MiB
```

```
0.851 ms [1]
2.386 ms [162761]
3.920 ms [13423]
5.455 ms [1382]
6.989 ms [297]
8.524 ms [67]
10.058 ms [19]
11.593 ms [4]
13.127 ms [1]
14.662 ms [1]
16.196 ms [5]
```

10.00%	in	1.4333	ms
25.00%	in	1.4679	ms
50.00%	in	1.5116	ms
75.00%	in	1.6109	ms
90.00%	in	2.0064	ms
95.00%	in	2.9872	ms
99.00%	in	3.9197	ms
99.90%	in	6.1850	ms
99.99%	in	9.1293	ms

DNS+dialup: 0.4255 ms, 0.2872 ms, 0.5384 ms
DNS-lookup: 0.1721 ms, 0.0070 ms, 0.3988 ms

[200] 177961 responses

Part 4 - Stress Test (No caching)

```
oha -z 30s -c 200 http://localhost:3000/page1.html
```

```
Success rate: 100.00%
Total:      30002.9344 ms
Slowest:    721.4011 ms
```

```
Total data: 148.08 MiB
Size/request: 873 B
Size/sec: 4.94 MiB
```

Response time histogram:

20.020	ms	[1]	
90.158	ms	[177773]	
160.296	ms	[10]	
230.434	ms	[11]	
300.572	ms	[13]	
370.711	ms	[10]	
440.849	ms	[11]	
510.987	ms	[9]	
581.125	ms	[10]	
651.263	ms	[9]	
721.401	ms	[8]	

Response time distribution:

10.00%	in	29.8013	ms
25.00%	in	31.2252	ms
50.00%	in	32.6732	ms
75.00%	in	34.8140	ms
90.00%	in	37.5157	ms
95.00%	in	40.5963	ms
99.00%	in	50.7841	ms
99.90%	in	77.5339	ms
99.99%	in	579.9252	ms

Details (average, fastest, slowest):

```
DNS+dialup: 12.3355 ms, 0.1685 ms, 21.7581 ms
DNS-lookup: 0.0269 ms, 0.0012 ms, 0.7919 ms
```

Status code distribution:

[200] 177865 responses

```
oha -z 30s -c 400 http://localhost:3000/page1.html
```

Summary:

```
Success rate: 100.00%
Total: 30.0057 sec
Slowest: 2.0149 sec
Fastest: 0.0465 sec
Average: 0.0697 sec
Requests/sec: 5747.1054
```

Total data: 143.24 MiB
Size/request: 873 B
Size/sec: 4.77 MiB

0.046	sec	[1]	
0.243	sec	[171922]	
0.440	sec	[16]	
0.637	sec	[15]	
0.834	sec	[15]	
1.031	sec	[14]	
1.228	sec	[14]	
1.424	sec	[14]	
1.621	sec	[12]	
1.818	sec	[12]	
2.015	sec	[12]	

10.00%	in 0.0629	sec
25.00%	in 0.0654	sec
50.00%	in 0.0680	sec
75.00%	in 0.0717	sec
90.00%	in 0.0770	sec
95.00%	in 0.0809	sec
99.00%	in 0.0912	sec
99.90%	in 0.1310	sec
99.99%	in 1.7349	sec

```
DNS+dialup: 0.0129 sec, 0.0003 sec, 0.0371 sec
DNS-lookup: 0.0000 sec, 0.0000 sec, 0.0006 sec
```

[200] 172047 responses

```
Summary:
  Success rate: 100.00%
  Total:      3.0011 10 sec
  Slowest:    1.7034 10 sec
  Fastest:    0.0047 10 sec
  Average:    0.0181 10 sec
  Requests/sec: 5527.1694

  Total data: 137.28 MiB
  Size/request: 873 B
  Size/sec: 4.57 MiB
```

```
0.005 10 sec [1] |
0.175 10 sec [164459] |
0.344 10 sec [63] |
0.514 10 sec [56] |
0.684 10 sec [53] |
```

0.100	ms	[1]
2.563	ms	[242438]
5.026	ms	[8801]
7.489	ms	[630]
9.951	ms	[140]
12.414	ms	[4]
14.877	ms	[1]
17.340	ms	[1]
19.802	ms	[1]
22.265	ms	[1]

Response time distribution:

10.00%	in	0.9933	ms
25.00%	in	1.0068	ms
50.00%	in	1.0333	ms
75.00%	in	1.0904	ms
90.00%	in	1.4034	ms
95.00%	in	2.3994	ms
99.00%	in	3.3475	ms
99.90%	in	6.8750	ms
99.99%	in	8.7983	ms

Details (average, fastest, slowest):

DNS+dialup: 0.4469 ms, 0.1914 ms, 0.6863 ms
DNS-lookup: 0.2882 ms, 0.0558 ms, 0.5218 ms

Status code distribution:

[200] 252019 responses

```
oha -z 30s -c 50 http://localhost:3000/page1.html
```

Summary :

```
Success rate: 100.00%
Total: 30000.8902 ms
Slowest: 246.8897 ms
Fastest: 0.1493 ms
Average: 6.0600 ms
Requests/sec: 8250.2219
```

```
Total data: 206.03 MiB
Size/request: 873 B
Size/sec: 6.87 MiB
```

Response time histogram:

0.149	ms	[1]	
24.823	ms	[247421]	
49.497	ms	[11]	
74.171	ms	[5]	
98.845	ms	[5]	
123.519	ms	[5]	
148.194	ms	[3]	
172.868	ms	[4]	
197.542	ms	[3]	
222.216	ms	[4]	
246.890	ms	[4]	

Response time distribution:

```
10.00% in 5.0942 ms
25.00% in 5.2001 ms
50.00% in 5.3954 ms
75.00% in 6.7622 ms
```


90.00%	in	7.7139	ms
95.00%	in	8.6535	ms
99.00%	in	11.5281	ms
99.90%	in	19.4486	ms
99.99%	in	91.2082	ms

Details (average, fastest, slowest):

DNS+ dialup: 0.2971 ms, 0.0961 ms, 0.7209 ms
DNS-lookup: 0.0733 ms, 0.0012 ms, 0.5982 ms

Status code distribution:

[200] 247466 responses

Error distribution:

```
[48] aborted due to deadline
```

Part 6 - Stress Test (caching)

```
oha -z 30s -c 200 http://localhost:3000/page1.html
```

Summary :

```
Success rate: 100.00%
Total: 30.0028 sec
Slowest: 2.0442 sec
Fastest: 0.0046 sec
Average: 0.0254 sec
Requests/sec: 7867.7265
```

```
Total data: 196.37 MiB
Size/request: 873 B
Size/sec: 6.55 MiB
```

Response time histogram:

0.005	sec	[1]	
0.209	sec	[235767]	
0.412	sec	[12]	
0.616	sec	[14]	
0.820	sec	[10]	
1.024	sec	[11]	
1.228	sec	[10]	
1.432	sec	[10]	
1.636	sec	[9]	
1.840	sec	[9]	
2.044	sec	[9]	

Response time distribution:

```
10.00% in 0.0221 sec
25.00% in 0.0232 sec
50.00% in 0.0240 sec
75.00% in 0.0258 sec
90.00% in 0.0301 sec
95.00% in 0.0327 sec
```

```
oha -z 30s -c 400 http://localhost:3000/page1.html
```

Summary :

```
Success rate: 100.00%
Total:      30.0042 sec
Slowest:    5.7871 sec
Fastest:    0.0015 sec
Average:    0.0509 sec
Requests/sec: 7850.4754
```

```
Total data: 195.78 MiB
Size/request: 873 B
Size/sec: 6.53 MiB
```

Response time histogram:

```
0.002 sec [1]
0.580 sec [235036]
1.159 sec [15]
1.737 sec [14]
2.316 sec [13]
2.894 sec [14]
3.473 sec [12]
4.051 sec [13]
4.630 sec [13]
5.209 sec [12]
5.787 sec [12]
```

Response time distribution:

10.00%	in	0.0436	sec
25.00%	in	0.0464	sec
50.00%	in	0.0482	sec
75.00%	in	0.0509	sec
90.00%	in	0.0562	sec
95.00%	in	0.0616	sec
99.00%	in	0.0789	sec
99.90%	in	0.1146	sec
99.99%	in	4.6471	sec

Details (average, fastest, slowest):

```
DNS+dialup: 0.0165 sec, 0.0003 sec, 0.0438 sec
DNS-lookup: 0.0000 sec, 0.0000 sec, 0.0018 sec
```

Status code distribution:

[200] 235155 responses

```
oha -z 30s -c 1000 http://localhost:3000/page1.html
```

```
Success rate: 100.00%
Total: 3.0014 10 sec
Slowest: 2.9840 10 sec
Fastest: 0.0026 10 sec
Average: 0.0106 10 sec
Requests/sec: 7955.3330
```

Response time histogram:

Response time distribution:

Details (average, fastest, slowest):

Status code distribution:

No Cache Stress test re-run

Summary :

```
Success rate: 100.00%
Total:      30.0045 sec
Slowest:    2.5887 sec
```

Fastest: 0.0450 sec
Average: 0.0677 sec
Requests/sec: 5906.3446

```
Total data: 147.22 MiB
Size/request: 873 B
Size/sec: 4.91 MiB
```

Response time histogram:

0.045	sec	[1]	
0.299	sec	[176671]	
0.554	sec	[23]	
0.808	sec	[17]	
1.062	sec	[18]	
1.317	sec	[20]	
1.571	sec	[18]	
1.826	sec	[16]	
2.080	sec	[16]	
2.334	sec	[16]	
2.589	sec	[15]	

Response time distribution:

10.00%	in	0.0608	sec
25.00%	in	0.0629	sec
50.00%	in	0.0655	sec
75.00%	in	0.0685	sec
90.00%	in	0.0745	sec
95.00%	in	0.0793	sec
99.00%	in	0.0902	sec
99.90%	in	0.1336	sec
99.99%	in	2.2871	sec

Details (average, fastest, slowest):

```
DNS+dialup: 0.0193 sec, 0.0003 sec, 0.0456 sec
DNS-lookup: 0.0000 sec, 0.0000 sec, 0.0008 sec
```

Status code distribution:

[200] 176831 responses